


# Network Programming

Eunji Lee  
([ejlee@ssu.ac.kr](mailto:ejlee@ssu.ac.kr))

# Putting it all Together: Anatomy of an Internet Connection



51213 is an ephemeral port  
allocated by the kernel

80 is a well-known port  
associated with Web servers

# Clients

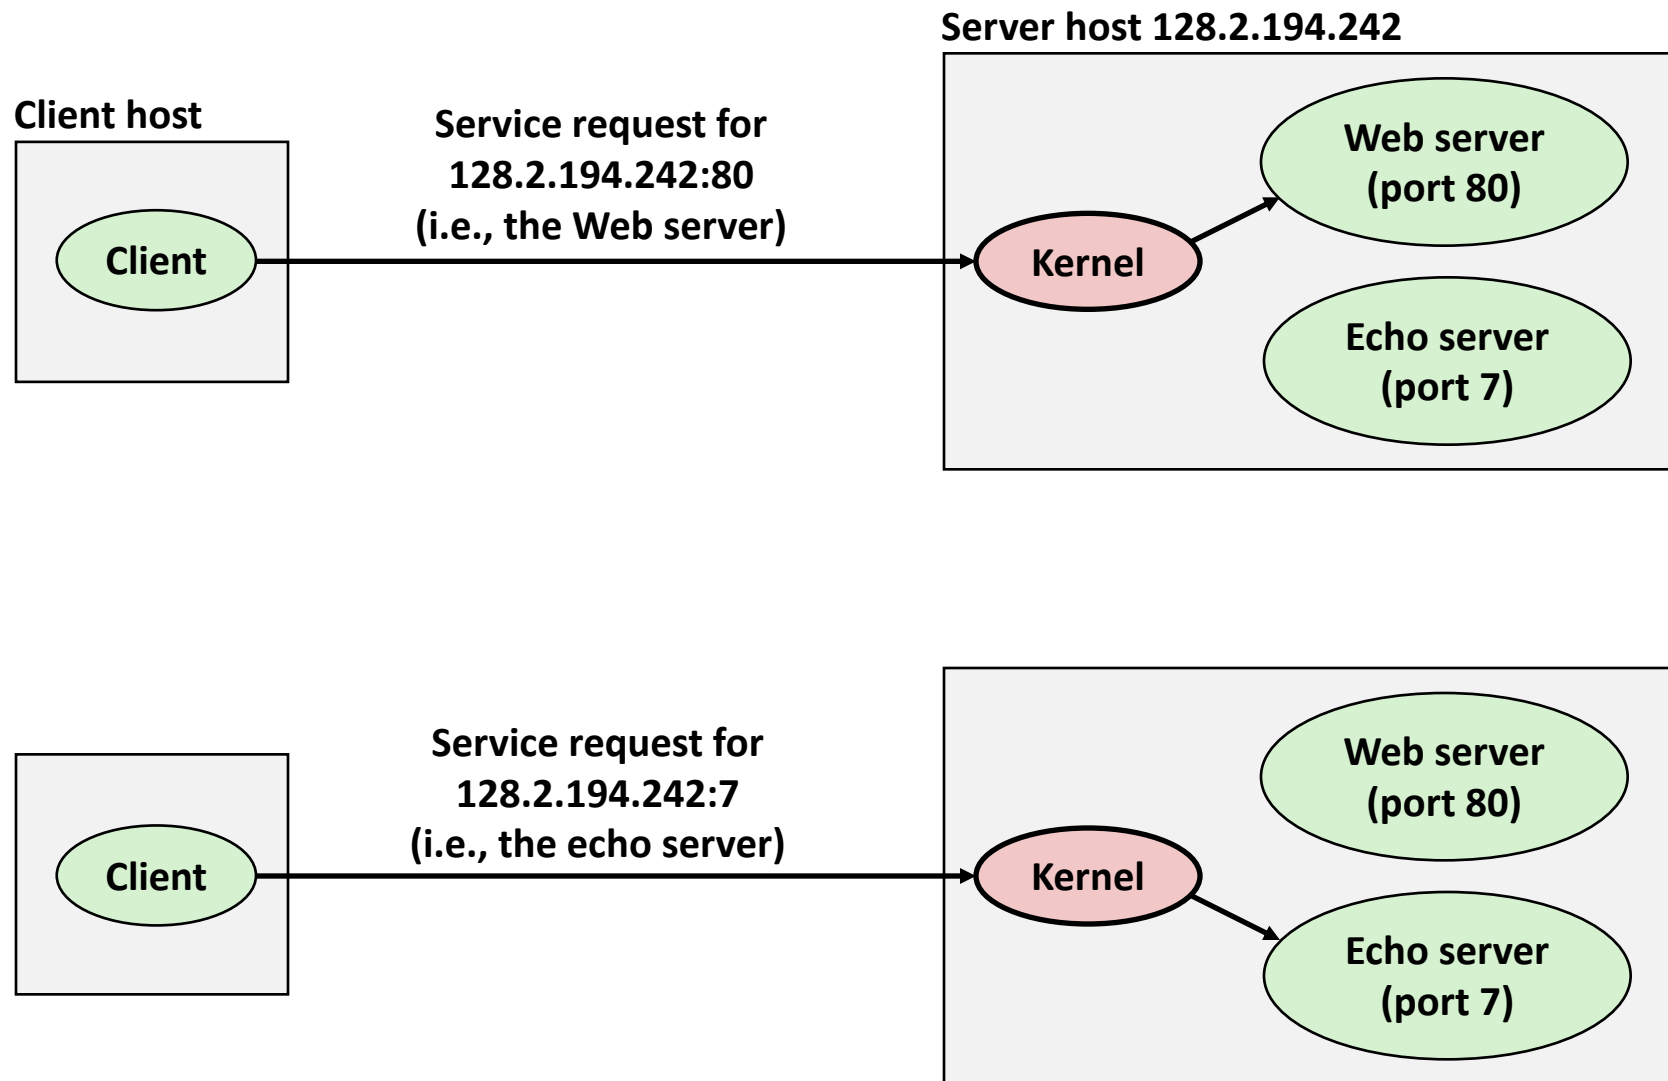
## ■ Examples of client programs

- Web browsers, `ftp`, `telnet`, `ssh`

## ■ How does a client find the server?

- The IP address in the server socket address identifies the host (more precisely, an adapter on the host)
- The (well-known) port in the server socket address identifies the service, and thus implicitly identifies the server process that performs that service.
- Examples of well know ports
  - Port 7: Echo server
  - Port 23: Telnet server
  - Port 25: Mail server
  - Port 80: Web server

# Using Ports to Identify Services



# Servers

- **Servers are long-running processes (daemons)**
  - Created at boot-time (typically) by the init process (process 1)
  - Run continuously until the machine is turned off
  
- **Each server waits for requests to arrive on a well-known port associated with a particular service**
  - Port 7: echo server
  - Port 23: telnet server
  - Port 25: mail server
  - Port 80: HTTP server
  
- **A machine that runs a server process is also often referred to as a “server”**

# Server Examples

## ■ Web server (port 80)

- Resource: files/compute cycles (CGI programs)
- Service: retrieves files and runs CGI programs on behalf of the client

## ■ FTP server (20, 21)

- Resource: files
- Service: stores and retrieve files

See `/etc/services` for a comprehensive list of the port mappings on a Linux machine

## ■ Telnet server (23)

- Resource: terminal
- Service: proxies a terminal on the server machine

## ■ Mail server (25)

- Resource: email “spool” file
- Service: stores mail messages in spool file

# Sockets Interface


- Created in the early 80's as part of the original Berkeley distribution of Unix that contained an early version of the Internet protocols
- Provides a user-level interface to the network
- Underlying basis for all Internet applications
- Based on client/server programming model

# Sockets

## ■ What is a socket?

- To the kernel, a socket is an endpoint of communication
- To an application, a socket is a file descriptor that lets the application read/write from/to the network
  - **Remember:** All Unix I/O devices, including networks, are modeled as files

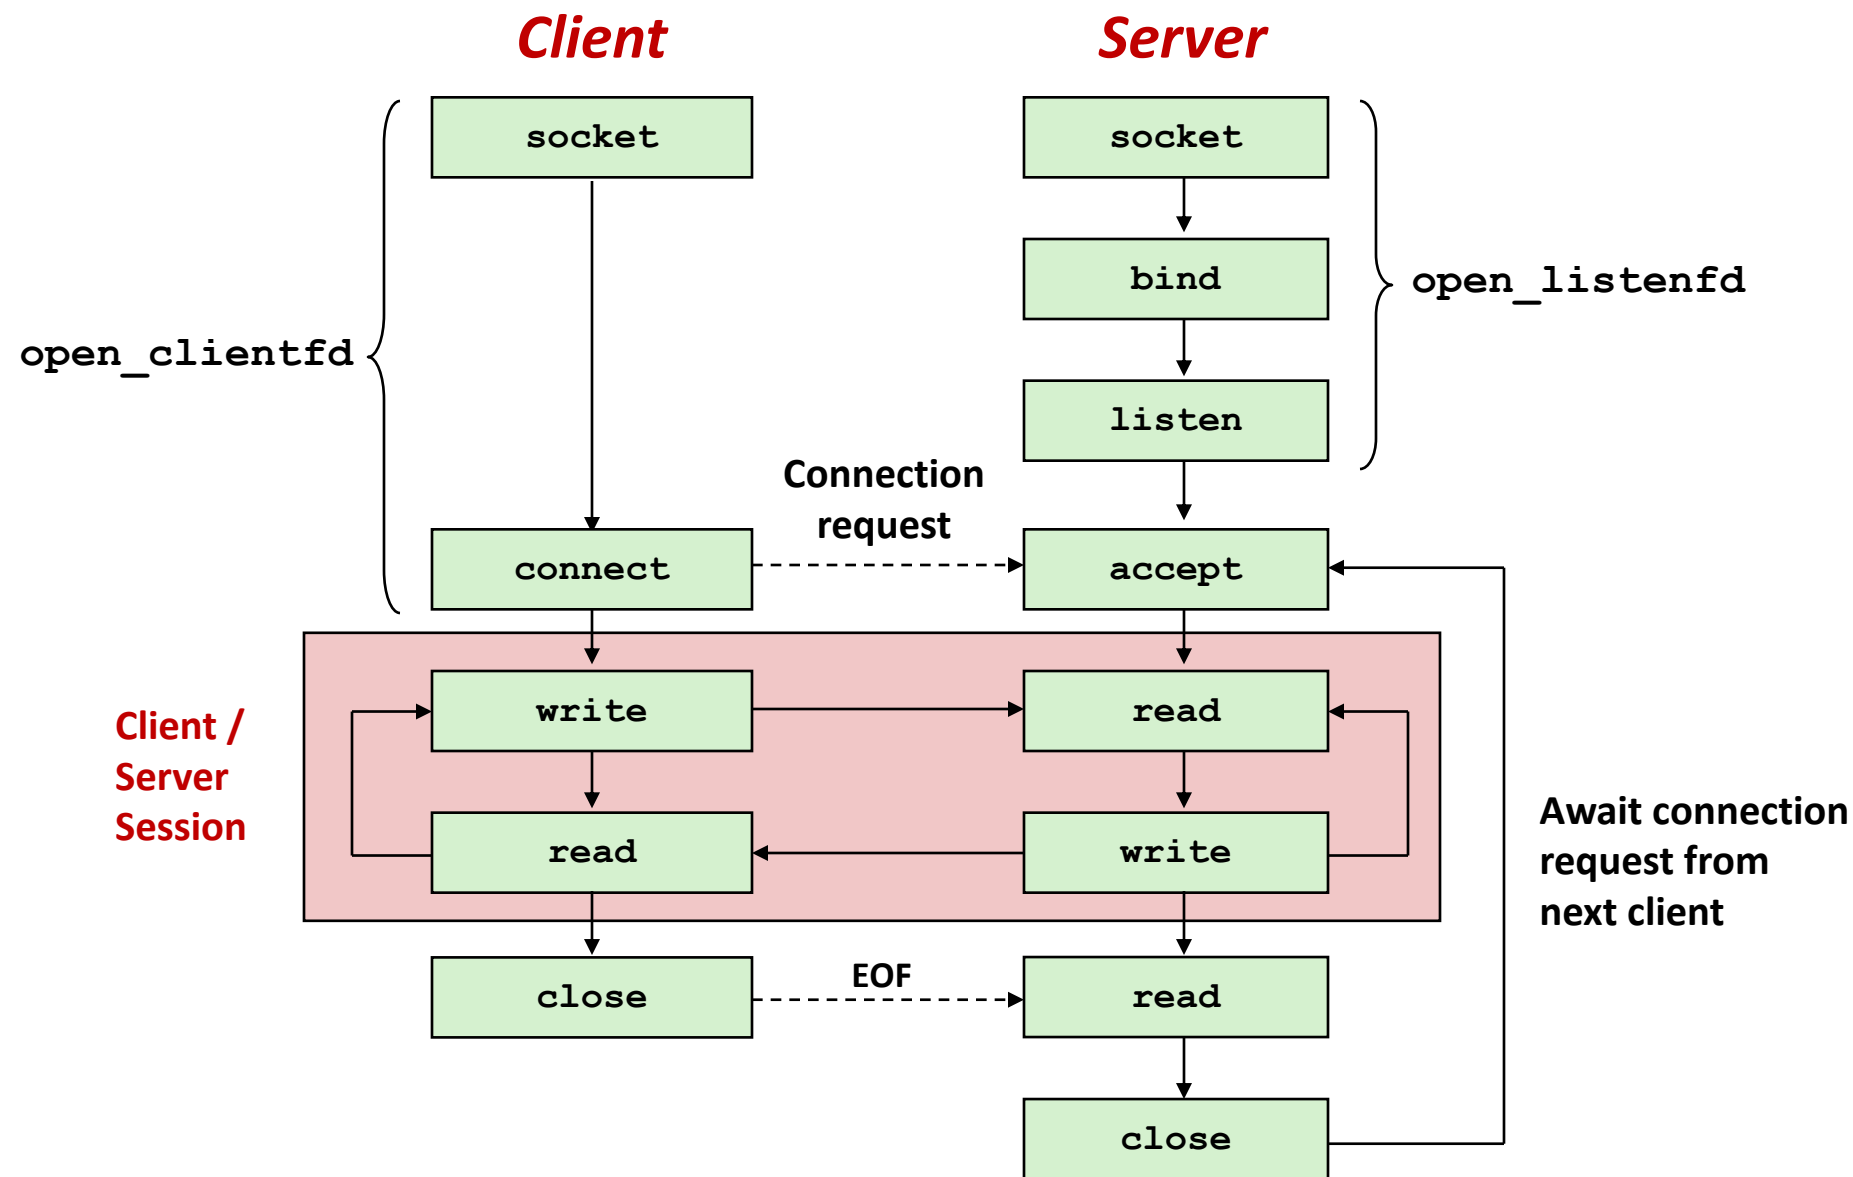
## ■ Clients and servers communicate with each other by reading from and writing to socket descriptors




- The main distinction between regular file I/O and socket I/O is how the application “opens” the socket descriptors



# Overview of the Sockets Interface



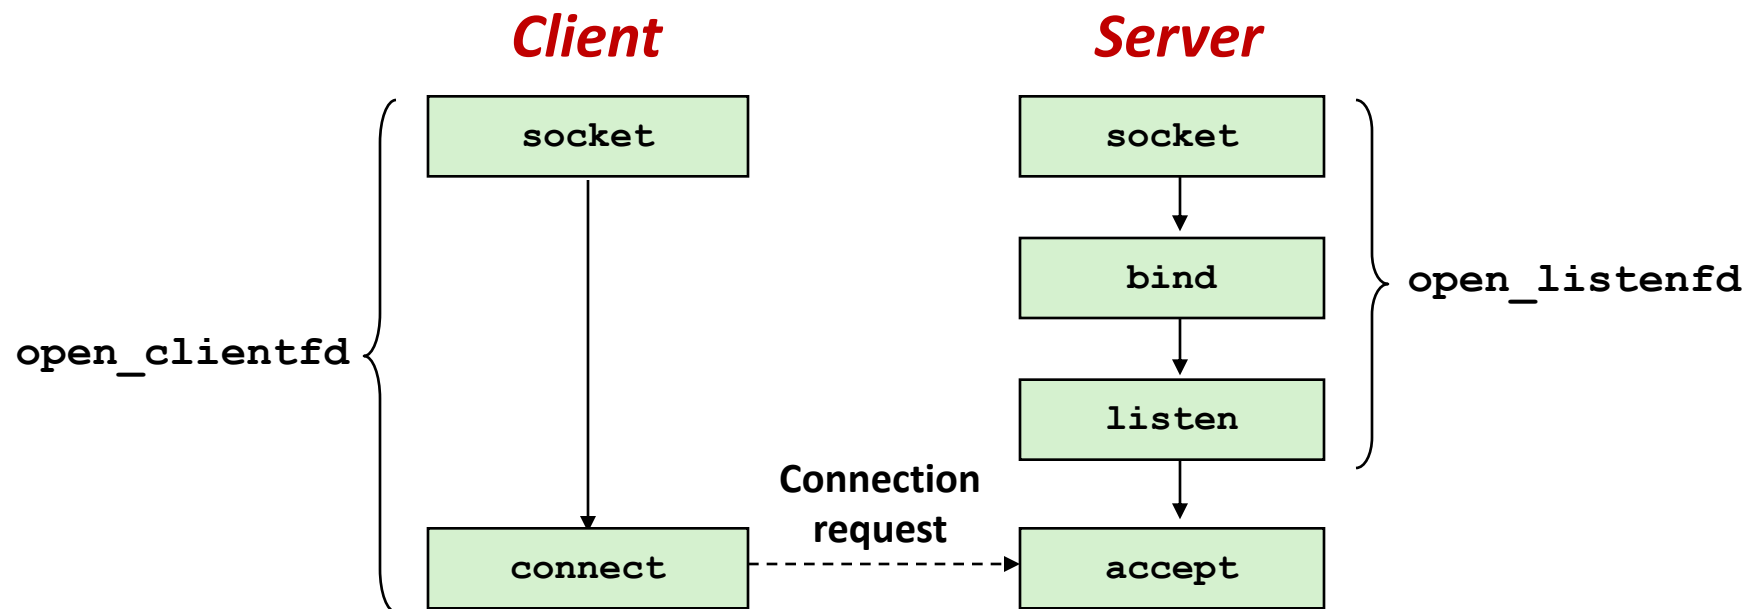
# Example: Echo Client / Server



```
(base) ejlee@datos-SYS-7049GP-TRT:~/lecture/sp_prof/net_lab$ ./client 127.0.0.1 8080
fd = 3
hello
hello
world
world
What a small world
What a small world
```

```
segmentation fault (core dumped)
(base) ejlee@datos-SYS-7049GP-TRT:~/lecture/sp_prof/net_lab$ ./server 8080
Waiting for request ...
Connected to 127.0.0.1
server received 6 bytes
server received 6 bytes
server received 19 bytes
```

# Overview of the Sockets Interface



## ■ Office Telephone Analogy for Server

- Socket: Buy a phone
- Bind: Tell the local administrator what number you want to use
- Listen: Plug the phone in
- Accept: Answer the phone when it rings

# Echo Server

```

1 #include<stdio.h>
2 #include<stdlib.h>
3 #include<string.h>
4 #include<unistd.h>
5 #include<sys/types.h>
6 #include<sys/socket.h>
7 #include<arpa/inet.h>
8
9 #define LISTENQ 1024
10 #define MAXLINE 1024
11
12 int main(int argc, char **argv)
13 {
14     int listenfd, connfd, port, clientlen;
15     struct sockaddr_in serveraddr, clientaddr;
16     char buf[MAXLINE];
17
18     // Create a socket
19     if ((listenfd = socket(AF_INET, SOCK_STREAM, 0)) < 0){
20         perror("socket");
21         exit(-1);
22     }
23
24     // Associate the socket with a pair of address and port number
25     port = atoi(argv[1]);
26     serveraddr.sin_family = AF_INET;
27     serveraddr.sin_addr.s_addr = htonl(INADDR_ANY);
28     serveraddr.sin_port = htons((unsigned short)port);
29
30     if (bind(listenfd, (struct sockaddr*)&serveraddr, sizeof(serveraddr)) == -1){
31         perror("Failed to bind");
32         return -1;
33     }
34
35     // Start waiting
36     if (listen(listenfd, LISTENQ) < 0)
37         return -1;
38
39     // Wait for connection request
40     printf("Waiting for request ... \n");
41
42     clientlen = sizeof(clientaddr);
43     if ((connfd = accept(listenfd, (struct sockaddr*)&clientaddr, &clientlen)) < 0){
44         perror("accept");
45         exit(-1);
46     }

```

```

47
48     char* caddrp = inet_ntoa(clientaddr.sin_addr);
49     printf("Connected to %s\n", caddrp);
50
51     int n;
52     while(1) {
53         n = read(connfd, buf, MAXLINE);
54         if (n == 0) break;
55         printf("server received %d bytes\n", n);
56         write(connfd, buf, n);
57     }
58     return 0;
59 }

```

# socket function

## ■ Create a socket descriptor

```
#include <sys/types.h>
#include <sys/socket.h>
```

```
int socket(int domain, int type, int protocol);
```

Returns: nonnegative descriptor if OK, -1 on error

## ■ Client

```
clientfd = Socket(AF_INET, SOCK_STREAM, 0);
```

- Just allocates & initializes some internal data structures
- AF\_INET: indicates that the socket is associated with Internet protocols
- SOCK\_STREAM: selects a reliable byte stream connection provided by TCP

# bind function

- Associate socket address in `my_addr` with the socket descriptor `sockfd`

```
#include <sys/socket.h>
```

```
int bind(int sockfd, struct sockaddr *my_addr, int addrlen);
```

Returns: 0 if OK, -1 on error

```
772
773     /* Listenfd will be an endpoint for all requests to port
774        on any IP address for this host */
775     bzero((char *) &serveraddr, sizeof(serveraddr));
776     serveraddr.sin_family = AF_INET;
777     serveraddr.sin_addr.s_addr = htonl(INADDR_ANY);
778     serveraddr.sin_port = htons((unsigned short)port);
779     if (bind(listenfd, (SA *)&serveraddr, sizeof(serveraddr)) < 0)
780         return -1;
```

# listen function

- **listen** indicates that this socket will accept connection (connect) requests from clients

```
#include <sys/socket.h>
```

```
int listen(int sockfd, int backlog);
```

Returns: 0 if OK, -1 on error

- **LISTENQ** is constant indicating how many pending requests allowed
- We're finally ready to enter the main server loop that accepts and processes client connection requests.

# accept function

- **accept ()** blocks waiting for a connection request

```
#include <sys/socket.h>
```

```
int accept(int listenfd, struct sockaddr *addr, int *addrlen);
```

Returns: nonnegative connected descriptor if OK, -1 on error


- **accept** returns a *connected descriptor* (connfd) with the same properties as the *listening descriptor* (listenfd)
  - Returns when the connection between client and server is created and ready for I/O transfers
  - All I/O with the client will be done via the connected socket
- **accept** also fills in client's IP address




# Echo Server: accept Illustrated



*1. Server blocks in `accept`, waiting for connection request on listening descriptor `listenfd`*



*2. Client makes connection request by calling and blocking in `connect`*



*3. Server returns `connfd` from `accept`. Client returns from `connect`. Connection is now established between `clientfd` and `connfd`*

# Connected vs. Listening Descriptors

## ■ Listening descriptor

- End point for client connection requests
- Created once and exists for lifetime of the server

## ■ Connected descriptor

- End point of the connection between client and server
- A new descriptor is created each time the server accepts a connection request from a client
- Exists only as long as it takes to service client

## ■ Why the distinction?

- Allows for concurrent servers that can communicate over many client connections simultaneously
  - E.g., Each time we receive a new request, we fork a child to handle the request

# Echo Client

```
1 #include<stdio.h>
2 #include<stdlib.h>
3 #include<string.h>
4 #include<unistd.h>
5 #include<arpa/inet.h>
6
7 #define MAXLINE 1024
8
9 int main(int argc, char** argv)
10 {
11     int clientfd, port;
12     struct sockaddr_in serveraddr;
13
14     char *host, buf[MAXLINE];
15
16     host = argv[1];
17     port = atoi(argv[2]);
18
19     if((clientfd = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
20         perror("socket");
21         exit(1);
22     }
23
24     serveraddr.sin_family = AF_INET;
25     serveraddr.sin_port = htons((unsigned short)port);
26     serveraddr.sin_addr.s_addr = inet_addr(host);
27
28     if(connect(clientfd, (struct sockaddr*)&serveraddr, sizeof(serveraddr)) < 0) {
29         perror("connect");
30         exit(1);
31     }
32
33     int rbytes, wbytes;
34     while(fgets(buf, MAXLINE, stdin) != NULL) {
35         wbytes = write(clientfd, buf, strlen(buf));
36         if(wbytes < strlen(buf))
37             printf("Failed to send message\n");
38         rbytes = read(clientfd, buf, MAXLINE);
39         fputs(buf, stdout);
40     }
41
42     close(clientfd);
43     return 0;
44 }
45
46
```

# connect function

- Establish an Internet connection with the server at socket address `serv_addr`

```
#include <sys/socket.h>
```

```
int connect(int sockfd, struct sockaddr *serv_addr, int addrlen);
```

Returns: 0 if OK, -1 on error

```

_____ sockaddr: socketbits.h (included by socket.h), sockaddr_in: netinet/in.h
/* Generic socket address structure (for connect, bind, and accept) */
struct sockaddr {
    unsigned short  sa_family; /* Protocol family */
    char            sa_data[14]; /* Address data. */
};

/* Internet-style socket address structure */
struct sockaddr_in {
    unsigned short  sin_family; /* Address family (always AF_INET) */
    unsigned short  sin_port;   /* Port number in network byte order */
    struct in_addr  sin_addr;    /* IP address in network byte order */
    unsigned char   sin_zero[8]; /* Pad to sizeof(struct sockaddr) */
};
_____ sockaddr: socketbits.h (included by socket.h), sockaddr_in: netinet/in.h

```

# For More Information

- **W. Richard Stevens, “Unix Network Programming: Networking APIs: Sockets and XTI”, Volume 1, Second Edition, Prentice Hall, 1998**
  - THE network programming bible
- **Unix Man Pages**
  - Good for detailed information about specific functions
- **Complete versions of the echo client and server are developed in the text**
  - Updated versions linked to course website
  - Feel free to use this code in your assignments

# (Backups) Execution of Echo Server / Client

## ■ Source codes

- <http://csapp.cs.cmu.edu/3e/code.html>


## ■ Client

```
lee — ejlee@oslab-SYS-1029U-TRT: ~/lecture/sp_prof/code/netp — ssh ◀ -bash — 92x38
[ejlee@oslab-SYS-1029U-TRT:~/lecture/sp_prof/code/netp$ ./echoclient
usage: ./echoclient <host> <port>
[ejlee@oslab-SYS-1029U-TRT:~/lecture/sp_prof/code/netp$ ./echoclient localhost 8080
Hi, there
Hi, there
Hello World
Hello World
```

## ■ Server

```
lee — ejlee@oslab-SYS-1029U-TRT: ~/lecture/sp_prof/code/netp — ssh ◀ -bash -
[ejlee@oslab-SYS-1029U-TRT:~/lecture/sp_prof/code/netp$ ./echoserveri 8080
server connected to localhost (127.0.0.1)
server received 10 bytes
server received 12 bytes
```

# (Backups) Example: Echo Client / Server



```
$ ./echoclient localhost 8080  
Hi, there  
Hi, there  
Hello World  
Hello World
```

```
$ ./echoserveri 8080  
server connected to localhost  
(127.0.0.1)  
server received 10 bytes  
server received 12 bytes
```

# (Backups) Echo Client: Main Routine

```
1 /*
2  * echoclient.c - An echo client
3  */
4 /* $begin echoclientmain */
5 #include "csapp.h"
6
7 int main(int argc, char **argv)
8 {
9     int clientfd, port;
10    char *host, buf[MAXLINE];
11    rio_t rio;
12
13    if (argc != 3) {
14        fprintf(stderr, "usage: %s <host> <port>\n", argv[0]);
15        exit(0);
16    }
17    host = argv[1];
18    port = atoi(argv[2]);
19
20    clientfd = Open_clientfd(host, port);
21    Rio_readinitb(&rio, clientfd);
22
23    while (Fgets(buf, MAXLINE, stdin) != NULL) {
24        Rio_writen(clientfd, buf, strlen(buf));
25        Rio_readlineb(&rio, buf, MAXLINE);
26        Fputs(buf, stdout);
27    }
28    Close(clientfd); //line:netp:echoclient:close
29    exit(0);
30 }
31 /* $end echoclientmain */
```



# (Backups) Echo Client: open\_clientfd

```
721 *****/
722 /*
723 * open_clientfd - open connection to server at <hostname, port>
724 *   and return a socket descriptor ready for reading and writing.
725 *   Returns -1 and sets errno on Unix error.
726 *   Returns -2 and sets h_errno on DNS (gethostbyname) error.
727 */
728 /* $begin open_clientfd */
729 int open_clientfd(char *hostname, int port)
730 {
731     int clientfd;
732     struct hostent *hp;
733     struct sockaddr_in serveraddr;
734
735     if ((clientfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)
736         return -1; /* check errno for cause of error */
737
738     /* Fill in the server's IP address and port */
739     if ((hp = gethostbyname(hostname)) == NULL)
740         return -2; /* check h_errno for cause of error */
741     bzero((char *) &serveraddr, sizeof(serveraddr));
742     serveraddr.sin_family = AF_INET;
743     bcopy((char *)hp->h_addr_list[0],
744         (char *)&serveraddr.sin_addr.s_addr, hp->h_length);
745     serveraddr.sin_port = htons(port);
746
747     /* Establish a connection with the server */
748     if (connect(clientfd, (SA *) &serveraddr, sizeof(serveraddr)) < 0)
749         return -1;
750     return clientfd;
751 }
752 /* $end open_clientfd */
```

Create socket

Create address

Establish connection

# (Backups) accept function

```

1  /*
2  * echoserveri.c - An iterative echo server
3  */
4  /* $begin echoserverimain */
5  #include "csapp.h"
6
7  void echo(int connfd);
8
9  int main(int argc, char **argv)
10 {
11     int listenfd, connfd, port, clientlen;
12     struct sockaddr_in clientaddr;
13     struct hostent *hp;
14     char *haddrp;
15     if (argc != 2) {
16         fprintf(stderr, "usage: %s <port>\n", argv[0]);
17         exit(0);
18     }
19     port = atoi(argv[1]);
20
21     listenfd = Open_listenfd(port);
22     while (1) {
23         clientlen = sizeof(clientaddr);
24         connfd = Accept(listenfd, (SA *)&clientaddr, &clientlen);
25
26         /* determine the domain name and IP address of the client */
27         hp = Gethostbyaddr((const char *)&clientaddr.sin_addr.s_addr,
28             sizeof(clientaddr.sin_addr.s_addr), AF_INET);
29         haddrp = inet_ntoa(clientaddr.sin_addr);
30         printf("server connected to %s (%s)\n", hp->h_name, haddrp);
31
32         echo(connfd);
33         Close(connfd);
34     }
35     exit(0);
36 }
37 /* $end echoserverimain */

```

'echoserveri.c' 37L, 912C written

```
#include <sys/socket.h>
```

```
int accept(int listenfd, struct sockaddr *addr, int *addrlen);
```

Returns: nonnegative connected descriptor if OK, -1 on error

# (Backups) connect function

```
#include <sys/socket.h>

int connect(int sockfd, struct sockaddr *serv_addr, int addrlen);
Returns: 0 if OK, -1 on error
```

t <hostname, port>  
reading and writing.  
stbyname) error.

```
/* Generic socket address structure */
struct sockaddr {
    unsigned short  sa_family; /*
    char           sa_data[14]; /*
};
```

```
728 /* $begin open_clientfd */
729 int open_clientfd(char *hostname, int port)
730 {
731     int clientfd;
732     struct hostent *hp;
733     struct sockaddr_in serveraddr;

    if ((clientfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)
        return -1; /* check errno for cause of error */

    /* Fill in the server's IP address and port */
    if ((hp = gethostbyname(hostname)) == NULL)
        return -2; /* check h_errno for cause of error */
    bzero((char *) &serveraddr, sizeof(serveraddr));
    serveraddr.sin_family = AF_INET;
    bcopy((char *) hp->h_addr_list[0],
          (char *) &serveraddr.sin_addr.s_addr, hp->h_length);
    serveraddr.sin_port = htons(port);

    /* Establish a connection with the server */
    if (connect(clientfd, (SA *) &serveraddr, sizeof(serveraddr)) < 0)
        return -1;
    return clientfd;
751 }
752 /* $end open_clientfd */
```

```
struct sockaddr_in {
    unsigned short  sin_family;
    unsigned short  sin_port;
    struct in_addr  sin_addr;
    unsigned char   sin_zero[8];
};
```

# (Backups) Echo Server : Main Routine

```

1  /*
2  * echoserveri.c - An iterative echo server
3  */
4  /* $begin echoserverimain */
5  #include "csapp.h"
6
7  void echo(int connfd);
8
9  int main(int argc, char **argv)
10 {
11     int listenfd, connfd, port, clientlen;
12     struct sockaddr_in clientaddr;
13     struct hostent *hp;
14     char *haddrp;
15     if (argc != 2) {
16         fprintf(stderr, "usage: %s <port>\n", argv[0]);
17         exit(0);
18     }
19     port = atoi(argv[1]);
20
21     listenfd = Open_listenfd(port);
22     while (1) {
23         clientlen = sizeof(clientaddr);
24         connfd = Accept(listenfd, (SA *)&clientaddr, &clientlen);
25
26         /* determine the domain name and IP address of the client */
27         hp = Gethostbyaddr((const char *)&clientaddr.sin_addr.s_addr,
28             sizeof(clientaddr.sin_addr.s_addr), AF_INET);
29         haddrp = inet_ntoa(clientaddr.sin_addr);
30         printf("server connected to %s (%s)\n", hp->h_name, haddrp);
31
32         echo(connfd);
33         Close(connfd);
34     }
35     exit(0);
36 }
37 /* $end echoserverimain */

```

'echoserveri.c' 37L, 912C written

# (Backups) Echo Server: open\_listenfd

```

754 /*
755  * open_listenfd - open and return a listening socket on port
756  *     Returns -1 and sets errno on Unix error.
757  */
758 /* $begin open_listenfd */
759 int open_listenfd(int port)
760 {
761     int listenfd, optval=1;
762     struct sockaddr_in serveraddr;
763
764     /* Create a socket descriptor */
765     if ((listenfd = socket(AF_INET, SOCK_STREAM, 0)) < 0)
766         return -1;
767
768     /* Eliminates "Address already in use" error from bind. */
769     if (setsockopt(listenfd, SOL_SOCKET, SO_REUSEADDR,
770         (const void *)&optval , sizeof(int)) < 0)
771         return -1;
772
773     /* Listenfd will be an endpoint for all requests to port
774        on any IP address for this host */
775     bzero((char *) &serveraddr, sizeof(serveraddr));
776     serveraddr.sin_family = AF_INET;
777     serveraddr.sin_addr.s_addr = htonl(INADDR_ANY);
778     serveraddr.sin_port = htons((unsigned short)port);
779     if (bind(listenfd, (SA *)&serveraddr, sizeof(serveraddr)) < 0)
780         return -1;
781
782     /* Make it a listening socket ready to accept connection requests */
783     if (listen(listenfd, LISTENQ) < 0)
784         return -1;
785     return listenfd;
786 }
787 /* $end open_listenfd */

```

Create socket  
 Create address  
 bind & listen

# (Backups) Echo Server: echo

- The server uses RIO to read and echo text lines

```
1  /*
2  * echo - read and echo text lines until client closes connection
3  */
4  /* $begin echo */
5  #include "csapp.h"
6
7  void echo(int connfd)
8  {
9      size_t n;
10     char buf[MAXLINE];
11     rio_t rio;
12
13     Rio_readinitb(&rio, connfd);
14     while((n = Rio_readlineb(&rio, buf, MAXLINE)) != 0) {
15         printf("server received %ld bytes\n", n);
16         Rio_writen(connfd, buf, n);
17     }
18 }
19 /* $end echo */
20
```